
Apache Hive

— CMM705 - Big Data
Programming
Lecture 6 —

Overview

- Hive
 - Background of Hive
 - Hive vs Pig
 - Hive Architecture
 - Limitation of Hive
 - Data types
 - Data Models
 - Partitioning and bucketing
- Lab session on Hive
 - Setting up Hadoop on single node
 - Setting up Hive
 - Running Hive

Background of Hive

- Started at Facebook
- Data was collected into Oracle DB by nightly cron jobs
- Grew from 10 of GBs in 2006 to 1 TB/day in 2007 and now it's 10x higher
 - > 950 Million Users
 - > 500 TB per day
 - > 70k queries per day
 - > 300m photos per day
- Users (employees) know SQL well

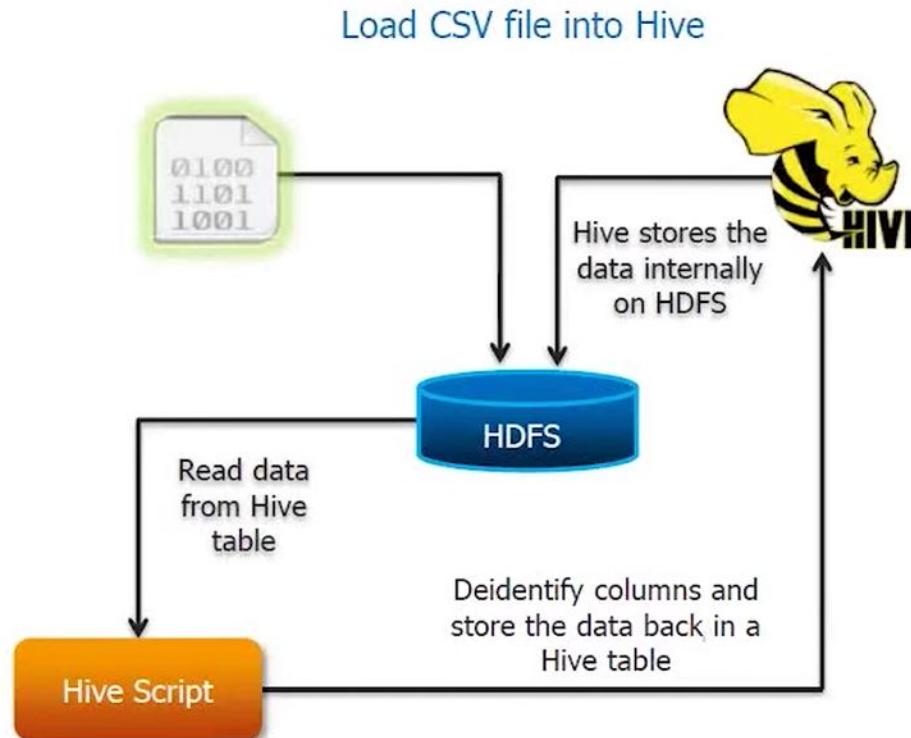
What is Hive ?

- A data warehousing package built on top of Hadoop
- Can be used for data analytics
- Targeted on users who are familiar with SQL
- Called as HiveQL
- No need to learn Java and Hadoop API :)
- Cannot process semi structured data.
 - Additional parsing (transformations) or schema inference needed
 - process semi-structured data by using UDFs (User Defined Functions) or by transforming data into a structured format
- Enable easy data ETL
 - Built-in Functions and Libraries

Hive

- Partitioning based on column values and Bucketing further segments to manageable chunks
- Schema flexibility (Schema on read) and evolution (Changing schema)
- Easy to plug-in custom mapreduce code
- JDBC/ODBC drivers are available
- HIVE tables can be defined directly on HDFS
- Extensible types, formats, functions & scripts
- Write once read many times (store years worth of data and analyze)
- Not a RDBMS or a database, just resembling a RDBMS for convenience

Hive (Managed Table)



Where to use Hive

- Log Processing
 - identify patterns, relationships, or trends within your data
- Customer facing BI - Top 10 users, etc
 - use ORDER BY or LIMIT clauses
- Data mining
 - identify patterns, relationships, or trends within your data
- Document Indexing
 - By creating tables that map to your document structure and indexing the table
- Predictive Modeling, Hypothesis Testing
 - Prepare and aggregate the data needed for your models
 - Exported to machine learning frameworks

Hive vs Pig ?

Pig

- Developed by Yahoo
- Procedural data flow

```
LOAD 'input_data' USING PigStorage(',')  
AS (field1:chararray, field2:int);  
FILTER A BY field2 > 100;  
GROUP B BY field1;  
FOREACH C GENERATE group, COUNT(B);
```

- Mostly used by Programmers and Researchers
- Implicit Schemas

Hive

- Developed by Facebook
- Declarative SQL

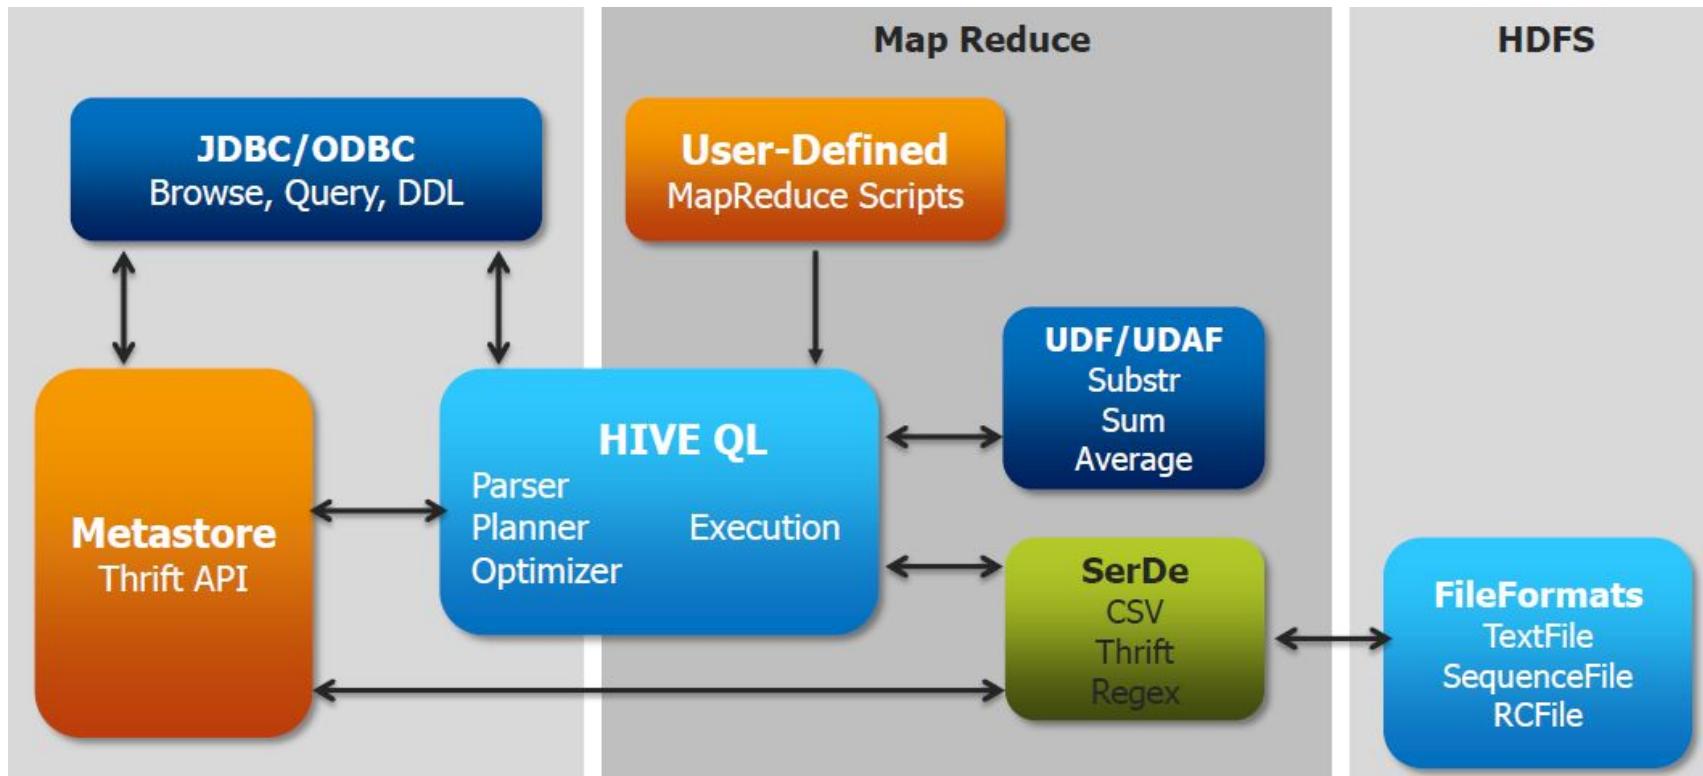
```
SELECT field1, COUNT(*) FROM input_data  
WHERE field2 > 100 GROUP BY field1;
```

- Mostly used by Analysts generating reports daily
- Provided partitions

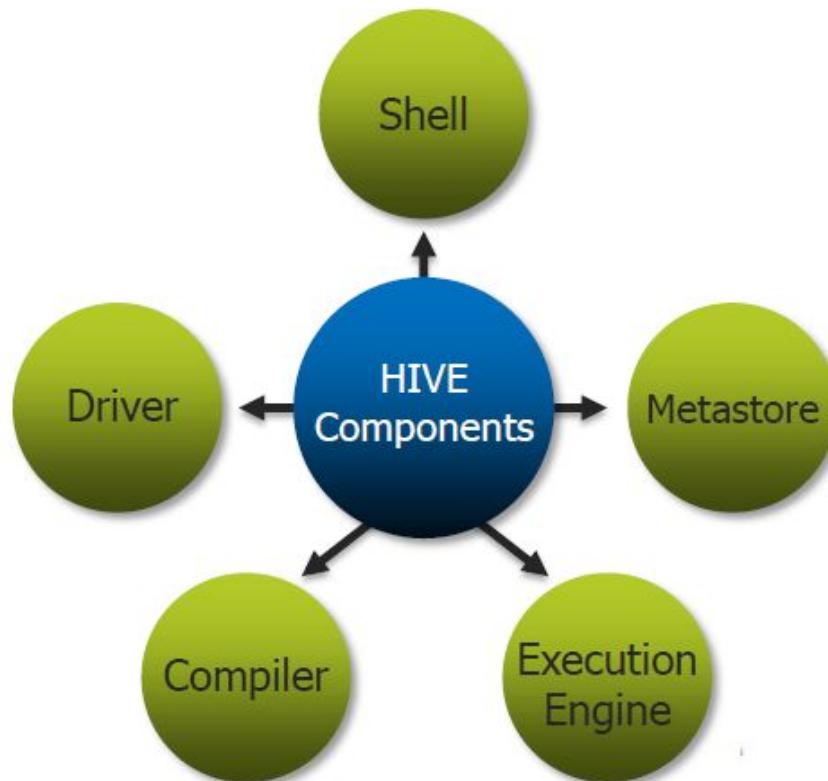
Hive vs Pig ? ...

Features	Hive	Pig
Language	SQL-like	PigLatin
Schemas/Types	Yes (explicit)	Yes (implicit)
Partitions	Yes	No
Server	Optional (Thrift)	No
User Defined Functions (UDF)	Yes (Java)	Yes (Java)
Custom Serializer/Deserializer	Yes	Yes
DFS Direct Access	Yes (implicit)	Yes (explicit)
Join/Order/Sort	Yes	Yes
Shell	Yes	Yes
Streaming	Yes	Yes
Web Interface	Yes	No
JDBC/ODBC	Yes (limited)	No

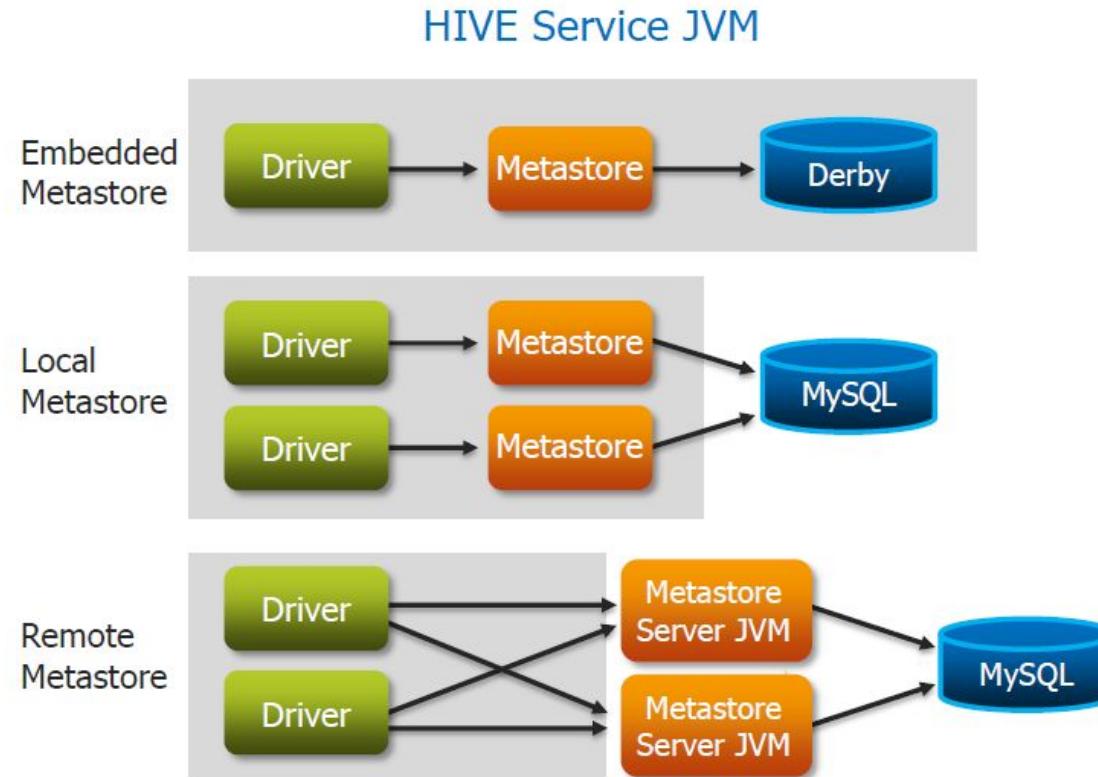
Hive Architecture



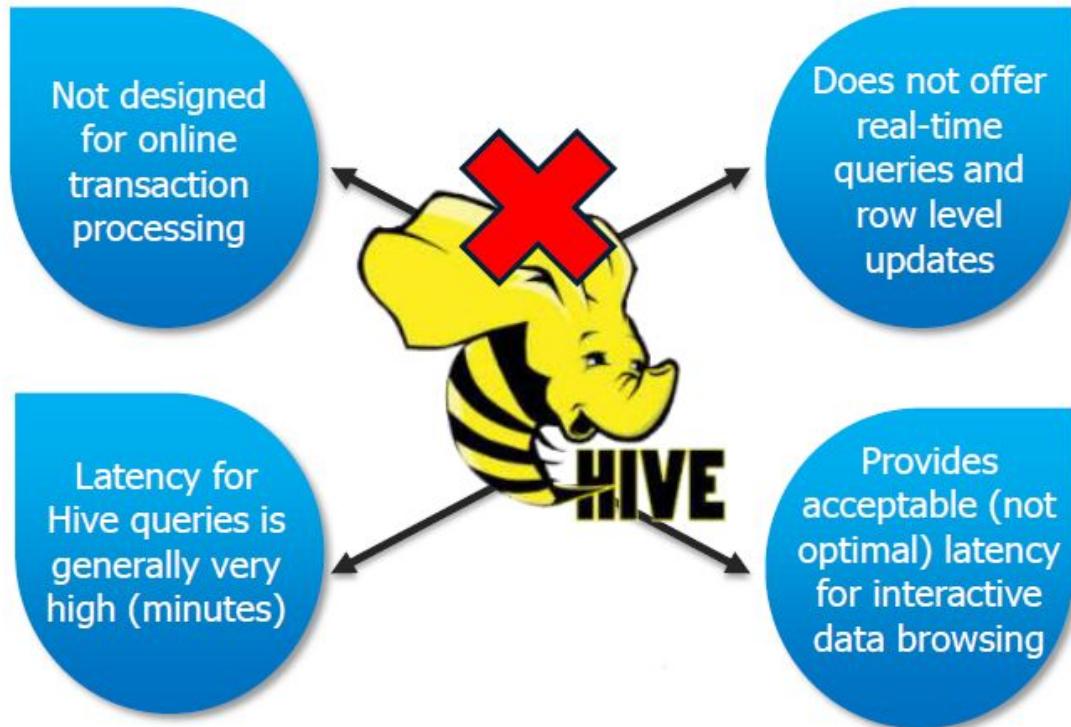
Components of Hive



Metastore

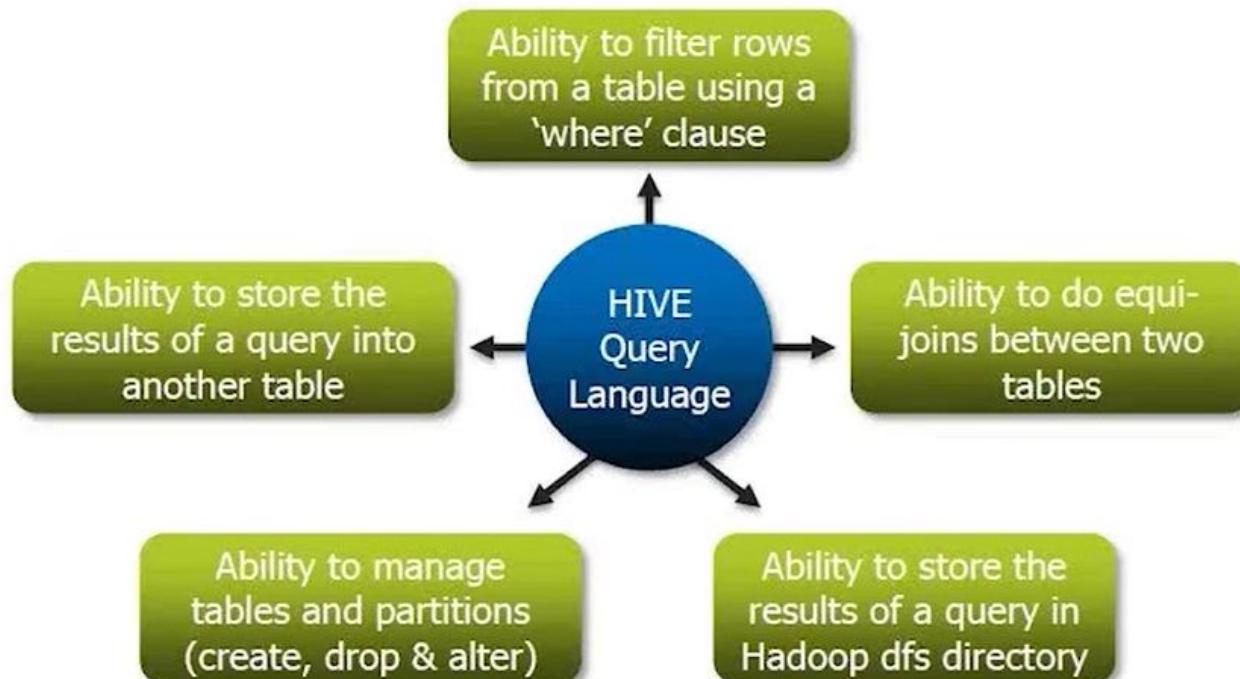


Limitation of Hive



Ability of HiveQL

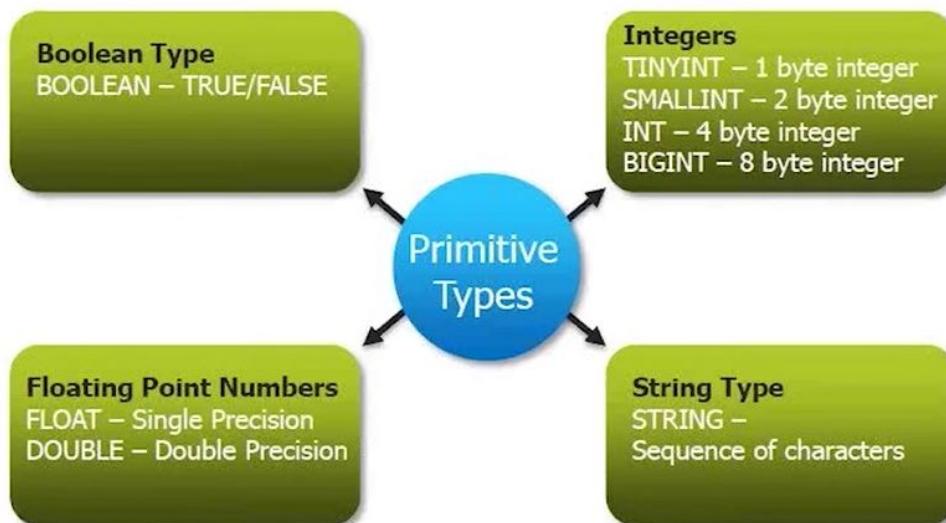
Hive Query Language provides the basic SQL-like operations



Data Handling in Hive

- Hive does not verify the data when it's loaded.
- It only verify the data when query is issued - Schema on Read
- No updates and transactions
 - Hive tables are generally append-only; once data is written to a table, you cannot update individual rows or cells within that table as you would in a traditional database
 - Did not support ACID (Atomicity, Consistency, Isolation, Durability) transactions
 - From Hive 0.14, Hive introduced ACID transactions with **limited transaction support**, specifically for insert, update, and delete operations

Data Types Supported



Array, Union

+

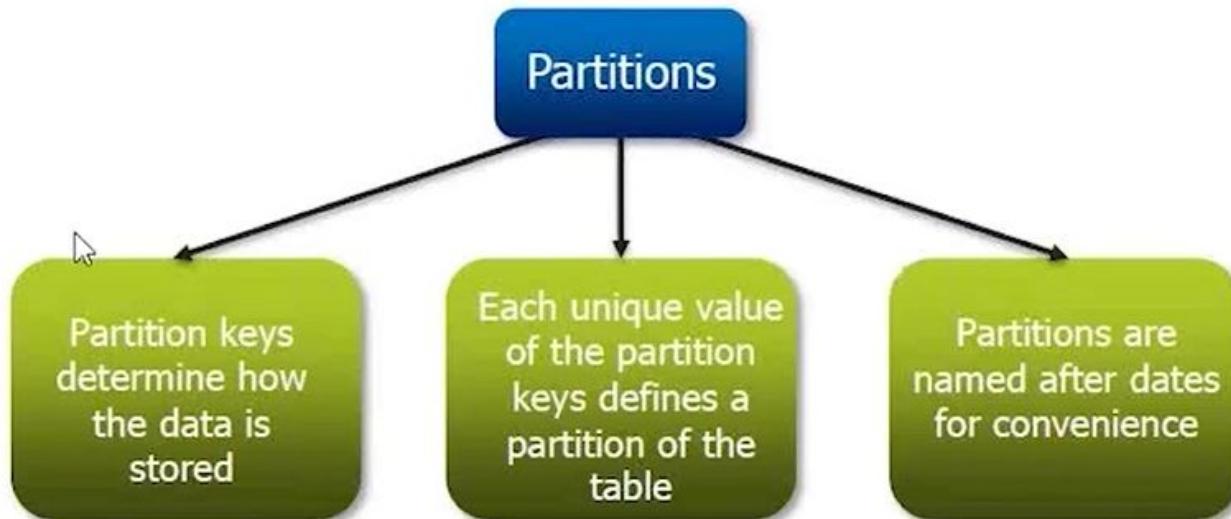
Struct, Map

Hive Data Models

- Database
 - Namespaces
- Tables
 - Schemas in namespaces
- Partitions
 - How data is stored in HDFS
 - Grouping data based on some column
- Buckets or Clusters
 - Partitions divided further into buckets based on some other column
 - Used for data sampling

Partitioning

Partition means dividing a table into coarse grained parts based on the value of a partition column such as a date. This make it faster to do queries on slices of the data.



Lab Session

Starting the Docker Image

- Clone the below git repository

```
git clone https://github.com/ramindu-msc/iit  
git pull origin main
```

Or if you have the repository created already

```
cd lab5/hadoop-hive-dockercompose/
```

- Change mapreduce-design-intro/hadoop-dockercompose/docker-compose.yaml's `/home/iitgcpuser/iit/lab5/hadoop-hive-dockercompose/resources` to your repository cloned path
- And start the containers with the following command(also by replacing the path)
`sudo docker compose -f
/home/iitgcpuser/iit/lab5/hadoop-hive-dockercompose/docker-compose-new.yaml up -d`
- Check the datanode is connected to namenode and fully started
`sudo docker logs -f hadoop-hive-dockercompose-datanode-1`

Running Hive Commands

- In case of an error as follows

```
x Container hive-metastore-init          service "hive-metastore-init" didn't
complete successfully: exit 1
```

- Run the following commands to get to hive shell and create database

```
sudo docker compose -f
/home/iitgcpuser/iit/lab5/hadoop-hive-dockercompose/docker-compose.yaml down -v

sudo docker volume rm $(sudo docker volume ls -q)
```

- Restart the containers

```
sudo docker compose -f
/home/iitgcpuser/iit/lab5/hadoop-hive-dockercompose/docker-compose.yaml up -d
```

Running Hive Commands

- Navigate to Run the docker image and run the following commands

```
sudo docker exec -it hive-server bash
```

- Navigate to Run the docker image and run the following commands

```
hive
create database telecom;
show databases;
create database telecom_backup comment 'holds backup data';
describe database extended telecom_backup;
```

Creating a Managed Table

- Open a separate terminal for namenode

```
docker exec -it namenode bash
```

```
hdfs dfs -ls /user/hive/warehouse/telecom.db/recharge
```

- Run the following command to create a table in hive shell

```
use telecom;
```

```
create table recharge( cell_no int, city string, name string, price float) row format
```

```
delimited fields terminated by ',' ;
```

```
describe extended recharge;
```

- In the terminal for namenode

```
hdfs dfs -ls /user/hive/warehouse/telecom.db/recharge
```

Creating a Managed Table..

- Run the following commands insert data, select data

```
INSERT INTO recharge (cell_no,city,name,price) VALUES (999090,"sl","fernando",30.0);  
SELECT * FROM recharge;
```

- Load data from local disk

```
LOAD DATA LOCAL INPATH '/opt/recharge2.input' INTO TABLE recharge;  
SELECT * FROM recharge;
```

- Run the following commands in namenode terminal

```
hdfs dfs -ls /user/hive/warehouse/telecom.db/recharge  
hdfs dfs -cat _____
```

Creating a External Table

- In the terminal for namenode

```
hdfs dfs -mkdir -p /user/hive/warehouse/telecom.db/recharge2
```

```
hdfs dfs -put /opt/recharge2.input /user/hive/warehouse/telecom.db/recharge2
```

```
hdfs dfs -cat /user/hive/warehouse/telecom.db/recharge2/recharge2.input
```

- Run the following commands in namenode shell

```
create external table recharge_external(cell_no int, city string, name string, price float)
```

```
row format delimited fields terminated by ',' LOCATION
```

```
'hdfs://namenode:8020/user/hive/warehouse/telecom.db/recharge2';
```

```
SELECT * FROM recharge_external;
```

Creating a External Table..

- In the terminal for namenode

```
echo -e "\n11436,s1,de silva,100" | hdfs dfs -appendToFile -  
/user/hive/warehouse/telecom.db/recharge2/recharge2.input
```

- In the terminal for namenode

```
echo "11436,s1,ramindu,100" > new_data.txt  
hdfs dfs -put -f new_data.txt /user/hive/warehouse/telecom.db/recharge2/recharge2.input
```

- Run the following commands in namenode shell

```
SELECT * FROM recharge_external;
```

Partitioning and Bucketing

```
use telecom;  
create table rechargeP (  
    cell_no int,  
    name string,  
    price float)  
partitioned by (city string)  
clustered by (name) into 10 buckets  
    row format delimited fields terminated by ','  
    stored as textfile;
```

- Same command in single line

```
create table rechargeP ( cell_no int, name string, price float) partitioned by (city string)  
clustered by (name) into 10 buckets row format delimited fields terminated by ',' stored as  
textfile; LOCATION 'hdfs://namenode:8020/user/hive/warehouse/telecom.db/partitioned';
```

Partitioning and Bucketing ...

- Set properties

```
set hive.exec.dynamic.partition.mode=nonstrict;  
set hive.exec.dynamic.partition=true;
```

- Load data to the partitioned table from recharge_external:

```
INSERT OVERWRITE TABLE rechargeP PARTITION (city) SELECT r2.cell_no, r2.name,  
r2.price, r2.city FROM recharge_external r2;
```

- See the partitions

```
SHOW PARTITIONS rechargeP;
```

Writing to Local Files

- Insert into local file

```
INSERT OVERWRITE LOCAL DIRECTORY '/temp/h_result' ROW FORMAT DELIMITED FIELDS  
TERMINATED BY ',' select * from recharge;
```

```
Exit;
```

- Check the local file

```
cat /temp/h_result/000000_0
```

What we covered

- Hive
 - Background of Hive
 - Hive vs Pig
 - Hive Architecture
 - Limitation of Hive
 - Data types
 - Data Models
 - Partitioning and bucketing
- Lab session on Hive
 - Setting up Hadoop on single node
 - Setting up Hive
 - Running Hive

Resources

- http://www.edureka.co/blog/hive-data-models/?utm_source=youtube&utm_medium=referral&utm_campaign=hive-tut1
- http://www.edureka.co/blog/pig-vs-hive/?utm_source=youtube&utm_medium=referral&utm_campaign=hive-tut1
- <https://cwiki.apache.org/confluence/display/Hive/GettingStarted#GettingStarted-InstallationandConfiguration>
- <https://www.youtube.com/watch?v=MoKW5eY5yVY>
- <https://www.youtube.com/watch?v=tKNGB5IZPFE>
- <http://hadoop.apache.org/docs/r2.7.2/hadoop-project-dist/hadoop-common/SingleCluster.html>